

Student Name: Adnan Ali, Student ID: 20020536

```
In [154]: import numpy as np
import pandas as pd
import wbgapi as wb
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.cluster as cluster
import scipy.optimize as opt
import itertools as iter

import warnings
warnings.filterwarnings("ignore")
```

```
In [155]: country_codes = ["GBR", "AUS", "USA", "DEU", "FRA", "JPN", "KOR", "CAN", "CHE", "TUR", "ESP"]

country_names = {
    "GBR" : "United Kingdom",
    "AUS" : "Australia",
    "USA" : "United States",
    "DEU" : "Germany",
    "FRA" : "France",
    "JPN" : "Japan",
    "KOR" : "South Korea",
    "CAN" : "Canada",
    "CHE" : "Switzerland",
    "TUR" : "Turkey",
    "ESP" : "Spain",
    "ITA" : "Italy"
}

#Agricultural Land (% of land area)
#Crop production
#Food production
indicators=['AG.LND.AGRI.ZS', 'AG.PRD.CROP.XD', 'AG.PRD.FOOD.XD']
```

```

In [156]: # This function retrieves data from world bank website and saves it into a file
def bank_data(data):
    print(wb.series.info(indicators))

    years = np.array(range(1970,2019))

    #MultiIndex Columns
    columns=pd.MultiIndex.from_arrays([[ 'Agricultural_land', 'Crop_production',
                                         'Food_production'],
                                       ['', '', '']])

    #MultiIndex Rows
    eco_index = pd.MultiIndex.from_product([data, years],
                                           names=[ 'Country', 'Year'])

    df_production = pd.DataFrame()

    for all_data in data:
        df_cli = wb.data.DataFrame(indicators, all_data, time=range(1970,2019)).s
        #use concat hierarchical indexing
        df_production = pd.concat([df_production, df_cli.T], ignore_index=True)

    df_production = df_production.set_index(eco_index)

    # Checking for null values in dataframe
    display("Checking null values in DataFrame")
    display(df_production.isnull().sum())

    return df_production

```

```
In [157]: final_df = bank_data(country_codes)
final_df
```

```
id          value
-----
AG.LND.AGRI.ZS  Agricultural land (% of land area)
AG.PRD.CROP.XD  Crop production index (2014-2016 = 100)
AG.PRD.FOOD.XD  Food production index (2014-2016 = 100)
3 elements

'Checking null values in DataFrame'

Agricultural_land    0
Crop_production      0
Food_production      0
dtype: int64
```

```
Out[157]:
```

		Agricultural_land	Crop_production	Food_production
--	--	-------------------	-----------------	-----------------

Country	Year			
	1970	77.910966	70.529999	77.970001
	1971	77.886165	73.989998	79.190002
GBR	1972	77.526557	71.080002	80.269997
	1973	77.278552	73.870003	81.239998
	1974	77.038813	76.220001	84.769997
...
	2014	44.747399	95.400002	95.580002
	2015	44.009655	103.059998	102.550003
ITA	2016	43.557490	101.540001	101.879997
	2017	43.607024	96.940002	98.070000
	2018	41.665267	95.150002	98.910004

588 rows × 3 columns

```
In [158]: final_df.groupby('Country').describe().T
```

Out[158]:

	Country	AUS	CAN	CHE	DEU	ESP	FR
Agricultural_land	count	49.000000	49.000000	49.000000	49.000000	49.000000	49.000000
	mean	58.259455	6.793882	40.238786	50.523189	59.564984	55.600600
	std	6.316012	0.185615	1.261064	2.341447	3.683630	2.372540
	min	44.539926	6.435684	38.214254	47.641536	52.408210	52.341700
	25%	51.853677	6.757391	39.142647	48.594861	57.342980	53.533600
	50%	60.383088	6.833125	40.142680	49.640453	60.345587	55.162100
	75%	63.883472	6.893133	41.524995	52.690975	62.449445	57.793900
	max	65.181261	7.044824	42.347197	54.486867	65.396775	59.578200
Crop_production	count	49.000000	49.000000	49.000000	49.000000	49.000000	49.000000
	mean	67.154286	62.104082	113.120000	90.664285	81.226122	92.835900
	std	23.693385	21.087347	13.459657	7.381608	16.670734	8.510100
	min	31.360001	30.680000	92.169998	74.239998	50.980000	72.019900
	25%	49.119999	46.400002	103.059998	84.610001	68.720001	89.279900
	50%	59.869999	58.900002	111.300003	91.489998	81.290001	94.150000
	75%	88.239998	72.870003	122.400002	95.510002	94.529999	99.510000
	max	121.239998	108.019997	152.089996	107.610001	121.690002	105.989900
Food_production	count	49.000000	49.000000	49.000000	49.000000	49.000000	49.000000
	mean	71.736123	69.242858	98.539388	93.981837	78.605510	97.674800
	std	17.264626	18.288463	5.127825	5.670311	17.397420	6.506800
	min	45.020000	41.340000	85.269997	83.339996	46.520000	82.750000
	25%	58.060001	55.779999	95.339996	89.309998	64.449997	95.849900
	50%	68.480003	65.889999	98.529999	94.139999	76.720001	99.239900
	75%	86.620003	82.000000	101.830002	97.500000	93.459999	102.029900
	max	104.040001	108.370003	108.660004	106.830002	117.860001	107.599900

```
In [159]: def norm(array):
min_val = np.min(array)
max_val = np.max(array)

scaled = (array-min_val) / (max_val-min_val)

return scaled


def norm_df(df):
    # iterate over all columns
    for col in df.columns[1:]: # excluding the first column
        df[col] = norm(df[col])


    return df
```

```
In [160]: # normalise result and inspect results
normalised_df = norm_df(final_df)
print(normalised_df.describe())
print()
```

	Agricultural_land	Crop_production	Food_production
count	588.000000	588.000000	588.000000
mean	44.253281	0.482929	0.575915
std	19.453966	0.198634	0.220691
min	6.435684	0.000000	0.000000
25%	34.683912	0.350404	0.411393
50%	49.325650	0.517750	0.650675
75%	57.848934	0.606004	0.734774
max	77.910966	1.000000	1.000000

```
In [161]: def makeplot(df, col1, col2):  
    plt.figure(figsize=(10.0,7.0))  
    plt.plot(df[col1], df[col2], "o", markersize=9)  
  
    plt.xlabel(col1)  
    plt.ylabel(col2)  
    plt.savefig(f'{col1}_dataplot.jpg')  
  
    plt.show()  
  
makeplot(normalised_df, "Agricultural_land", "Crop_production")  
makeplot(normalised_df, "Agricultural_land", "Food_production")
```





K Mean Clustering

```

In [162]: ##### kmeans set up the clusterer, 4 expected clusters
kmeans = cluster.KMeans(n_clusters=4)

# extract columns for fitting
df_fit = normalised_df[["Agricultural_land", "Crop_production"]].copy()
kmeans.fit(df_fit)

# extract labels and cluster centres
labels = kmeans.labels_
cen = kmeans.cluster_centers_

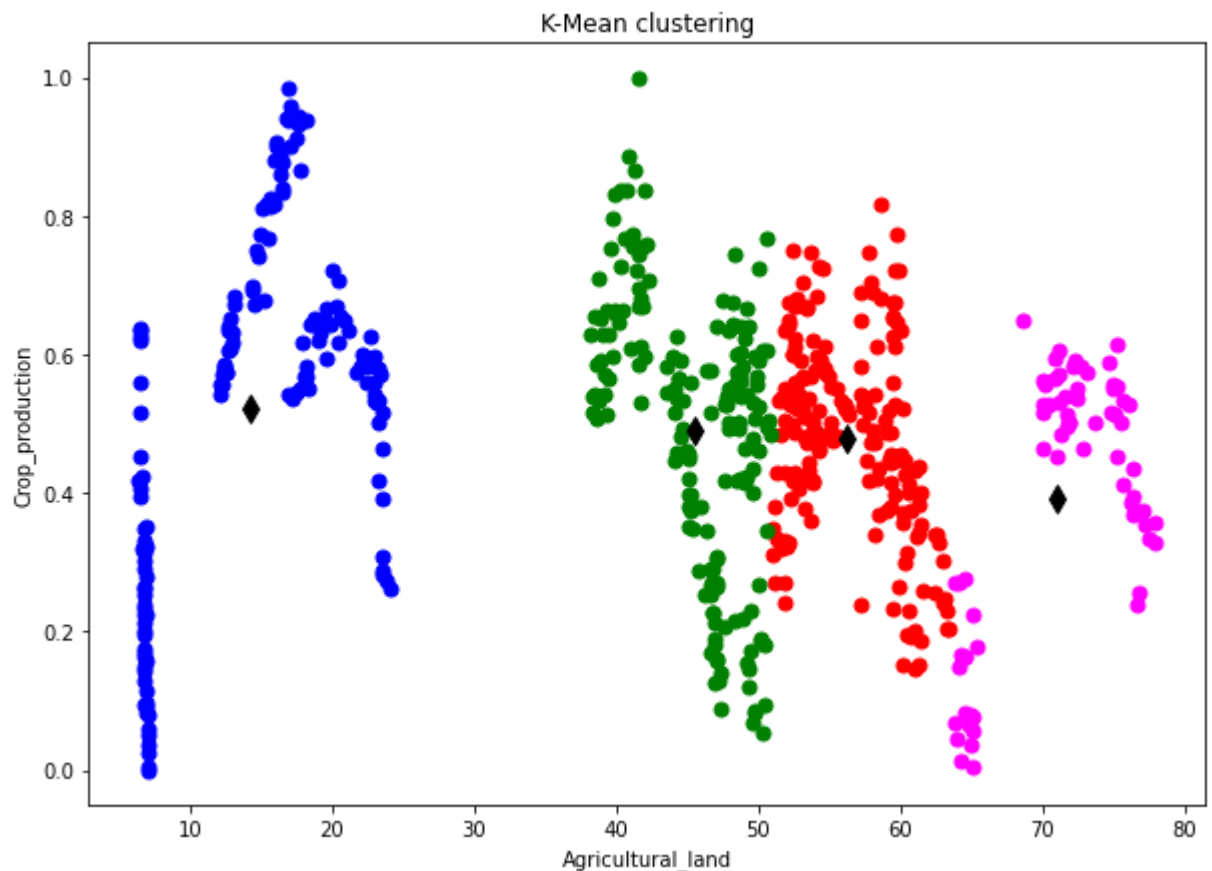
# plot using the labels to select colour
plt.figure(figsize=(10.0,7.0))

col = ["blue", "red", "green", "magenta"]
for l in range(4):      # loop over the different labels
    plt.plot(df_fit["Agricultural_land"][labels==l], df_fit["Crop_production"][labels==l],
             markersize=7, color=col[l])

# show cluster centres
for ic in range(4):
    xc, yc = cen[ic,:]
    plt.plot(xc, yc, "dk", markersize=10)

plt.xlabel("Agricultural_land")
plt.ylabel("Crop_production")
plt.title('K-Mean clustering')
plt.savefig('K-Mean cluster.jpg')
plt.show()

```



Agglomerative Clustering

```
In [163]: ##### setting up agglomerative clustering for 4 clusters
ac = cluster.AgglomerativeClustering(n_clusters=4)

df_fit = normalised_df[["Agricultural_land", "Crop_production"]].copy()

# carry out the fitting
ac.fit(df_fit)

labels = ac.labels_

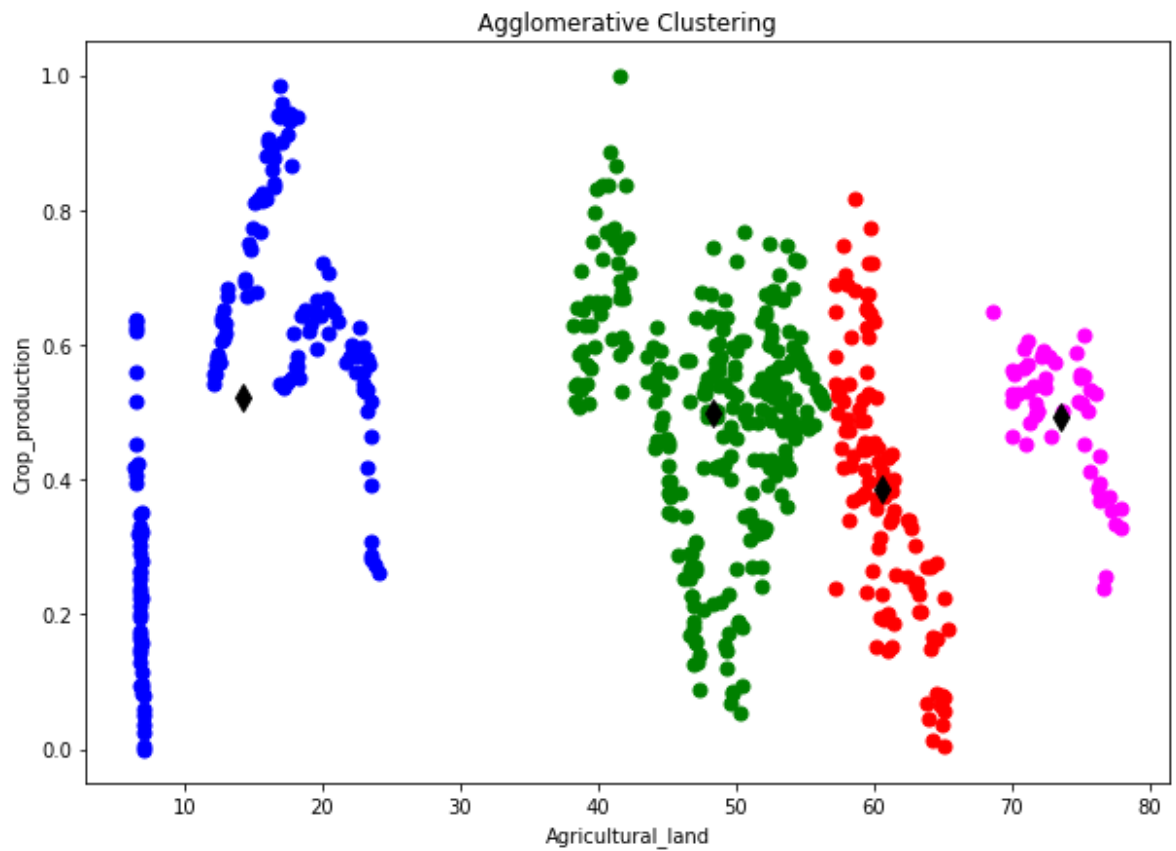
# The clusterer does not return cluster centres, but they are easily computed
xcen = []
ycen = []
for ic in range(4):
    xc = np.average(df_fit["Agricultural_land"][labels==ic])
    yc = np.average(df_fit["Crop_production"][labels==ic])
    xcen.append(xc)
    ycen.append(yc)

# plot using the labels to select colour
plt.figure(figsize=(10.0,7.0))

col = ["blue", "red", "green", "magenta"]
for l in range(0,4):      # loop over the different labels
    plt.plot(df_fit["Agricultural_land"][labels==l], df_fit["Crop_production"][labels==l],
             markersize=7, color=col[l])

# show cluster centres
for ic in range(4):
    plt.plot(xcen[ic], ycen[ic], "dk", markersize=10)

plt.xlabel("Agricultural_land")
plt.ylabel("Crop_production")
plt.title('Agglomerative Clustering')
plt.savefig('agglomerative cluster.jpg')
plt.show()
```



Data Fitting

```
In [164]: def exp_growth(t, scale, growth):  
            f = scale * np.exp(growth * (t-1990))  
            return f  
  
            def logistics(t, scale, growth, t0):  
                f = scale / (1.0 + np.exp(-growth * (t - t0)))  
                return f  
            # Define the exponential function and the Logistics functions for fitting.
```

```
In [165]: # taking one country from climate set
```

```
country_df = bank_data(['AUS'])
```

```
country_df
```

```
id          value
-----
AG.LND.AGRI.ZS  Agricultural land (% of land area)
AG.PRD.CROP.XD  Crop production index (2014-2016 = 100)
AG.PRD.FOOD.XD  Food production index (2014-2016 = 100)
3 elements
```

```
'Checking null values in DataFrame'
```

```
Agricultural_land    0
Crop_production       0
Food_production       0
dtype: int64
```

```
Out[165]:
```

		Agricultural_land	Crop_production	Food_production
Country	Year			
AUS	1970	64.229788	32.430000	45.020000
	1971	65.024298	35.200001	47.450001
	1972	65.060594	31.360001	48.590000
	1973	64.810929	38.509998	51.900002
	1974	65.131146	37.419998	47.040001
	1975	65.025708	40.529999	51.279999
	1976	65.181261	39.980000	54.880001
	1977	63.979537	36.320000	54.270000
	1978	64.577015	50.500000	62.419998
	1979	64.204470	48.619999	59.830002
	1980	64.516876	40.750000	52.590000
	1981	64.507374	50.650002	56.180000
	1982	63.883472	38.790001	51.369999
	1983	62.974500	60.080002	61.290001
	1984	63.334939	58.790001	58.060001
	1985	63.365268	55.450001	57.150002
	1986	60.989678	55.270000	59.099998
	1987	61.312367	49.200001	58.430000
	1988	61.446702	53.459999	60.970001
	1989	60.802624	54.110001	60.020000
	1990	60.455463	54.360001	62.720001
	1991	60.240032	49.119999	60.630001
	1992	60.659438	58.709999	65.760002

	Agricultural_land	Crop_production	Food_production
Country	Year		
1993	59.890528	62.919998	68.480003
1994	61.068040	48.279999	63.110001
1995	60.318915	66.940002	70.910004
1996	60.558687	80.680000	76.370003
1997	60.168569	75.779999	75.650002
1998	60.383088	82.419998	81.199997
1999	59.056923	89.760002	85.820000
2000	59.288104	84.760002	84.900002
2001	59.318173	86.040001	86.620003
2002	58.185700	88.239998	88.169998
2003	57.209429	59.740002	72.970001
2004	51.404657	95.300003	89.389999
2005	53.399373	88.370003	86.709999
2006	53.278575	92.639999	90.370003
2007	51.853677	59.869999	75.949997
2008	50.673366	72.800003	80.370003
2009	50.018484	86.779999	86.199997
2010	48.992099	86.940002	85.809998
2011	51.562813	95.650002	89.769997
2012	50.385431	104.099998	94.900002
2013	48.393710	98.379997	96.230003
2014	48.738399	99.040001	100.029999
2015	45.314424	98.620003	100.620003
2016	44.539926	102.330002	99.349998
2017	48.340618	121.239998	104.040001
2018	46.658095	93.330002	94.180000

```
In [166]: ctr=[]
          yr=[]
          for idx in country_df.index:
              ctr.append(idx[0])
              yr.append(idx[1])
          country_df.insert(0,"Country",ctr)
          country_df.insert(1,"Year",yr)
          country_df=country_df.reset_index(drop=True)
          country_df=country_df.fillna(country_df.mean())
          country_df["Year"] = country_df["Year"].astype(int)
```

```
In [167]: # fit exponential growth
popt, covar = opt.curve_fit(exp_growth, country_df['Year'], country_df['Crop_produ
```


```
In [168]: print("Fit parameter", popt)

# use *popt to pass on the fit parameters
country_df["crop_exp"] = exp_growth(country_df["Year"], *popt)

plt.figure(figsize=(10.0,7.0))
plt.plot(country_df["Year"], country_df["Crop_production"], label="data")
plt.plot(country_df["Year"], country_df["crop_exp"], label="fit")

plt.legend()
plt.title("First fit attempt")
plt.xlabel("Year")
plt.ylabel("Crop_production")
plt.savefig('exp_first attempt.jpg')
plt.show()
print()
```

Fit parameter [5.84383728e+01 2.25371029e-02]




```

In [169]: # decrease or increase exponential factor until rough agreement is reached
# growth of 0.02 gives a reasonable start value
popt = [0.32, 0.01]
country_df["crop_exp"] = exp_growth(country_df["Year"], *popt)

plt.figure(figsize=(10.0,7.0))
plt.plot(country_df["Year"], country_df["Crop_production"], label="data")
plt.plot(country_df["Year"], country_df["crop_exp"], label="fit")

plt.legend()
plt.xlabel("year")
plt.ylabel("Crop_production")
plt.title("Improved start value")
plt.savefig('exp_Improved start.jpg')
plt.show()

```



```

In [170]: # fit exponential growth
popt, covar = opt.curve_fit(exp_growth, country_df["Year"],
                             country_df["Crop_production"], p0=[0.32, 0.02])

# much better
print("Fit parameter", popt)


country_df["crop_exp"] = exp_growth(country_df["Year"], *popt)

plt.figure(figsize=(10.0,7.0))
plt.plot(country_df["Year"], country_df["Crop_production"], label="data")
plt.plot(country_df["Year"], country_df["crop_exp"], label="fit")

plt.legend()
plt.xlabel("Year")
plt.ylabel("Crop_production")
plt.title("Final fit exponential growth")
plt.savefig('exp_final fit.jpg')
plt.show()
print()

```

Fit parameter [5.84383657e+01 2.25371127e-02]





```

In [171]: # kept growth value from before
# increase scale factor and growth rate until rough fit
popt = [0.54, 0.02, 1990]
country_df["crop_log"] = logistics(country_df["Year"], *popt)

plt.figure(figsize=(10.0,7.0))
plt.plot(country_df["Year"], country_df["Crop_production"], label="data")
plt.plot(country_df["Year"], country_df["crop_log"], label="fit")

plt.legend()
plt.xlabel("year")
plt.ylabel("Crop_production")
plt.title("Improved start value")
plt.savefig('log_improved start.jpg')
plt.show()

```



```

In [187]: popt, covar = opt.curve_fit(logistics, country_df["Year"], country_df["Crop_production"],
                                     p0=(2e9, 0.05, 1990.0), maxfev=1200)
print("Fit parameter", popt)


country_df["crop_log"] = logistics(country_df["Year"], *popt)

plt.figure(figsize=(10.0,7.0))
plt.title("Logistics function")
plt.plot(country_df["Year"], country_df["Crop_production"], label="data")
plt.plot(country_df["Year"], country_df["crop_log"], label="fit")

plt.xlabel("Year")
plt.ylabel("Crop_production")
plt.savefig('log_function.jpg')
plt.show()

```

Fit parameter [1.63126671e+02 4.05412503e-02 2.00351105e+03]



```
In [188]: import itertools as iter

def err_ranges(x, func, param, sigma):
    # initiate arrays for lower and upper limits
    lower = func(x, *param)
    upper = lower

    uplow = [] # list to hold upper and lower limits for parameters
    for p,s in zip(param, sigma):
        pmin = p - s
        pmax = p + s
        uplow.append((pmin, pmax))

    pmix = list(iter.product(*uplow))

    for p in pmix:
        y = func(x, *p)
        lower = np.minimum(lower, y)
        upper = np.maximum(upper, y)


    return lower, upper
```

```
In [189]: # extract the sigmas from the diagonal of the covariance matrix
sigma = np.sqrt(np.diag(covar))

low, up = err_ranges(country_df["Year"], logistics, popt, sigma)

plt.figure(figsize=(10.0,7.0))
plt.title("Logistics function")
plt.plot(country_df["Year"], country_df["Crop_production"], label="data")
plt.plot(country_df["Year"], country_df["crop_log"], label="fit")

plt.fill_between(country_df["Year"], low, up, alpha=0.7)
plt.legend()
plt.xlabel("Year")
plt.ylabel("Crop_production")
plt.savefig('error range.jpg')
plt.show()
```



```
In [190]: print("Forecasted Crop_production")
low, up = err_ranges(2030, logistics, popt, sigma)
print("2030 between ", low, "and", up)
low, up = err_ranges(2040, logistics, popt, sigma)
print("2040 between ", low, "and", up)
low, up = err_ranges(2050, logistics, popt, sigma)
print("2050 between ", low, "and", up)
```

```
Forecasted Crop_production
2030 between 60.32698741111771 and 199.0963742721422
2040 between 67.83917716206273 and 206.7891088801082
2050 between 74.78712972215187 and 211.66156007700653
```

```
In [191]: print("Forcasted Crop_production")
low, up = err_ranges(2030, logistics, popt, sigma)
mean = (up+low) / 2.0
pm = (up-low) / 2.0
print("2030:", mean, "+/-", pm)

low, up = err_ranges(2040, logistics, popt, sigma)
mean = (up+low) / 2.0
pm = (up-low) / 2.0
print("2040:", mean, "+/-", pm)

low, up = err_ranges(2050, logistics, popt, sigma)
mean = (up+low) / 2.0
pm = (up-low) / 2.0
print("2050:", mean, "+/-", pm)
```

```
Forcasted Crop_production
2030: 129.71168084162997 +/- 69.38469343051224
2040: 137.31414302108544 +/- 69.47496585902273
2050: 143.2243448995792 +/- 68.43721517742733
```

```
In [ ]:
```